

HIGH PERFORMANCE COMPUTING for SCIENCE & ENGINEERING (HPCSE) I

HS 2022

TUTORIAL 01: BASICS

Athena Economides

Computational Science and Engineering Lab
ETH Zürich

21.09.2022

Outline

- I. Course organization: Moodle / website
- II. Git
- III. Makefiles
- IV. Euler Cluster: usage
- V. C++ refresher

I. Course Intro & Organization

TEACHING ASSISTANTS



- ▶ Athena: eceva@ethz.ch
- ▶ Michalis: michaich@ethz.ch
- ▶ Noah: noabauma@student.ethz.ch

COURSE WEBPAGE

<https://www.cse-lab.ethz.ch/teaching/hpcse-i-hs22/>

ASK QUESTIONS!

- ➔ Email us for administrative issues
- ➔ All other questions/ inquiries: Moodle
- ➔ Office hours: Tuesdays 9:00-10:00, Thursdays 12:00-13:00 (**CLT D11**)

HOMEWORKS

Not obligatory, can add *BONUS* to your final grade

final grade = $\max\{\text{exam grade}, (20\% \text{ hw grade} + 80\% \text{ exam grade})\}$

Submit homework 1-6 via the Moodle ETHZ app:

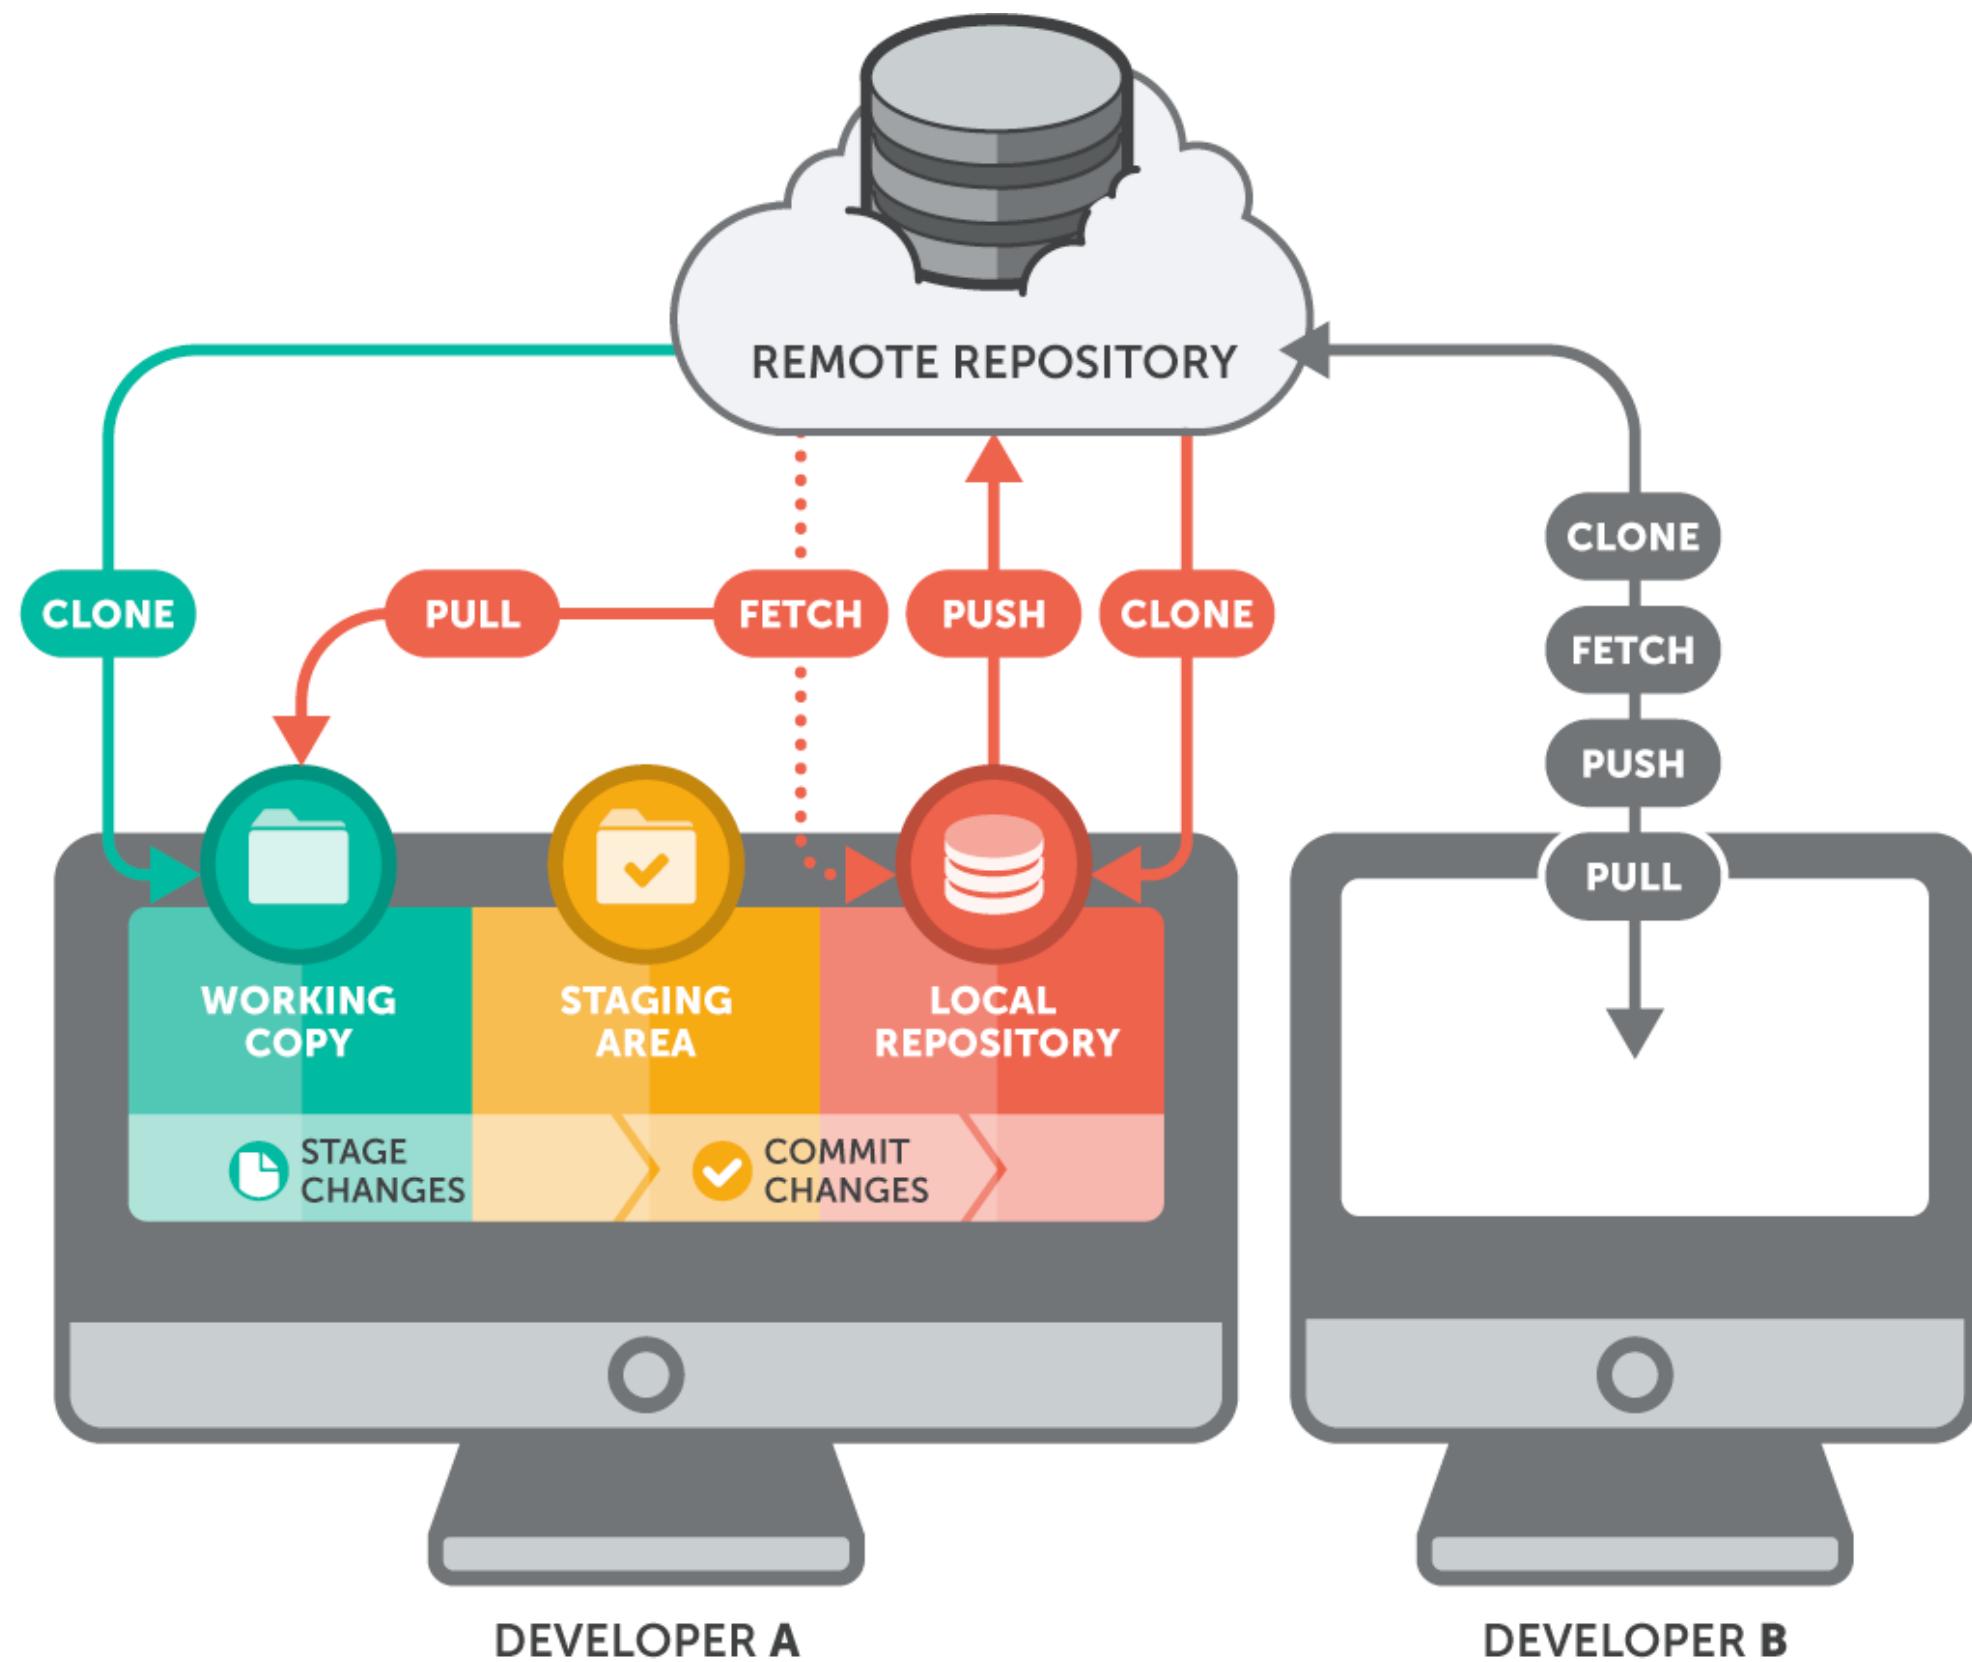
<https://moodle-app2.let.ethz.ch/course/view.php?id=18039>

II. Git repository

COURSE REPOSITORY

<https://gitlab.ethz.ch/hpcse-public-repos/hpcsei-fall-2022-lecture.git>

How does **Git** work?



Why **Git**?

- * One of the most popular version control systems (VCS) today!
- * One can save their work and work safely remotely and concurrently between multiple users
- * ETH provides access to gitlab.ethz.ch to all ETH members
- * Available on ETH clusters & Linux environments

HOW TO USE **Git**

<https://git-scm.com>

<https://git-scm.com/doc>

HOW TO INSTALL **Git**

<https://github.com/git-guides/install-git>

git clone <https://gitlab.ethz.ch/hpcse-public-repos/hpcsei-fall-2022-lecture.git>

git pull

III. Makefiles

CFLAGS: compiler
directives/ options

Possibly LDFLAGS:
list of link directives

Makefile 1 for factorial.cpp

compiler

targets

TAB!

```
CC=g++  
CFLAGS=-O3 -Wall -Wextra -Wpedantic
```

```
all: factorial
```

```
factorial.o: factorial.cpp
```

```
(CFLAGS)
```

```
factorial: factorial.o
```

```
$(CC) -o factorial factorial.o
```

```
clean:
```

```
rm -f *.o *~ factorial
```

```
.PHONY: all
```

```
.PHONY: clean
```

prerequisites

Compile commands

Convert .cpp code to object code

-c: generate object file

-o: put the compilation in the file factorial.o

Link the object code together into executable

COMMON MAKEFILE CFLAGS

- O3 : optimize generated code (00: not optimise code)
- g :compile with debug information
- std=c++11 use the c++11 standard language definition
- I\$(DIR) : include \$(DIR) containing auxiliary header files, other than /usr/include
- Wall : give verbose compiler warnings, include common warnings
- Wextra : enable extra warnings
- Wpedantic : turns off extensions and generates more warnings

USAGE

- make <target> : execute a specific target
- make clean : remove all object and executable files
- make : will execute the first target in the file

LDFLAGS EXAMPLES

- LDFLAGS=-lm : link the C math library
- LDFLAGS=-L../../lib -L<DIR>: for other <DIR> than /usr/lib

III. Makefiles

Dependency files
- all include files

- Suffix replacement rule:
- .o file depends upon the .cpp version of the file and the .hpp (DEPS) files
 - To generate the .o file, compile the .cpp file using the (CC) compiler.
 - -o \$@ : put the compilation in the file named on the left of : i.e. %.o
 - \$< : first item in dependencies list

```
CC=g++
CFLAGS=-O3 -Wall -Wextra -Wpedantic
OBJ=factorial.o
DEPS=factorial.hpp

all: factorial

%.o: %.cpp $(DEPS)
    $(CC) -c -o $@ $< $(CFLAGS)

factorial: $(OBJ)
    $(CC) -o $@ $^ $(CFLAGS)

clean:
    rm -f *.o *~ factorial

.PHONY: all
.PHONY: clean
```

Makefile 2 for factorial.cpp, where factorial(int n) is a function that calculates the factorial of an integer number.

Also in the directory, there exists a header file: factorial.hpp, with function declaration of factorial(int n)

- General compilation rule:
- \$@ \$^ : macros for the left and right sides of :, i.e. here: \$@=factorial and \$^=(OBJ)

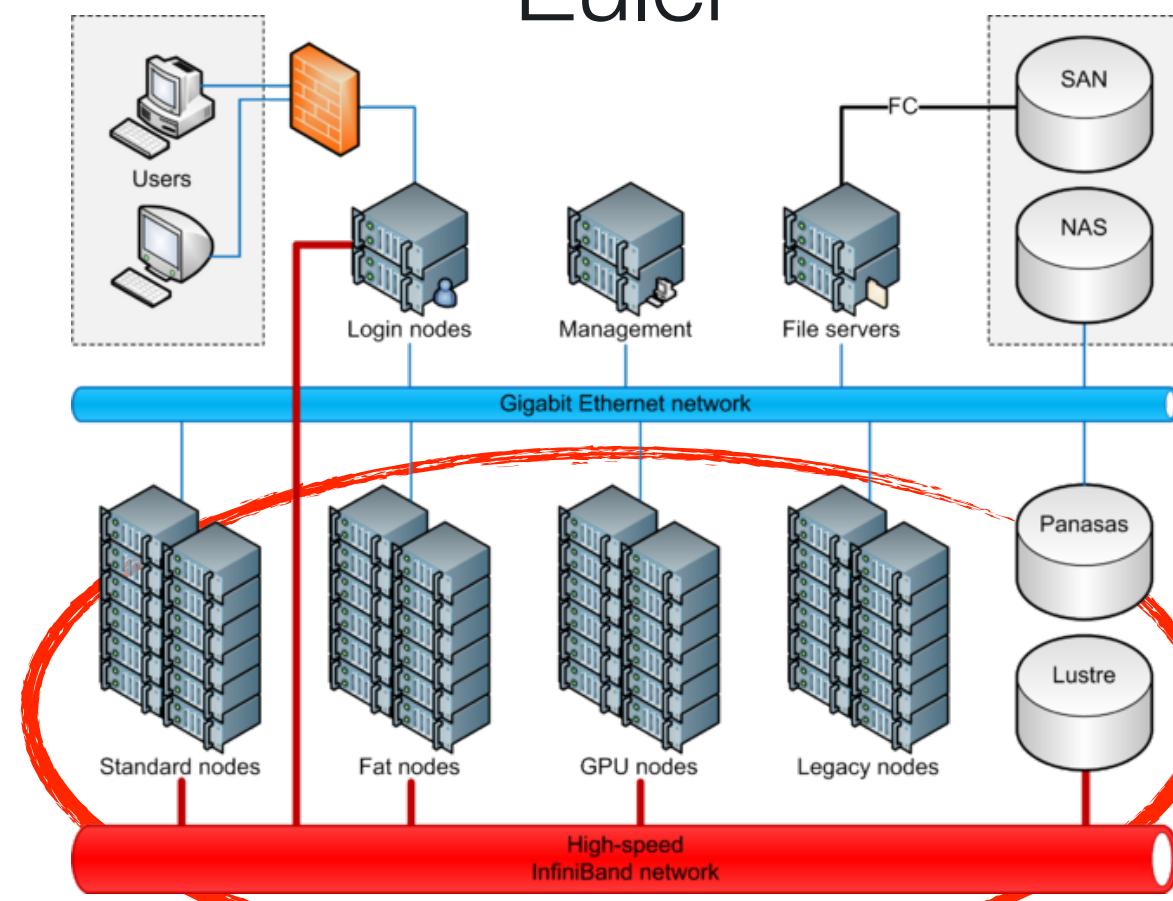
IV. Euler cluster : usage



DETAILS

<https://scicomp.ethz.ch/wiki/Euler>

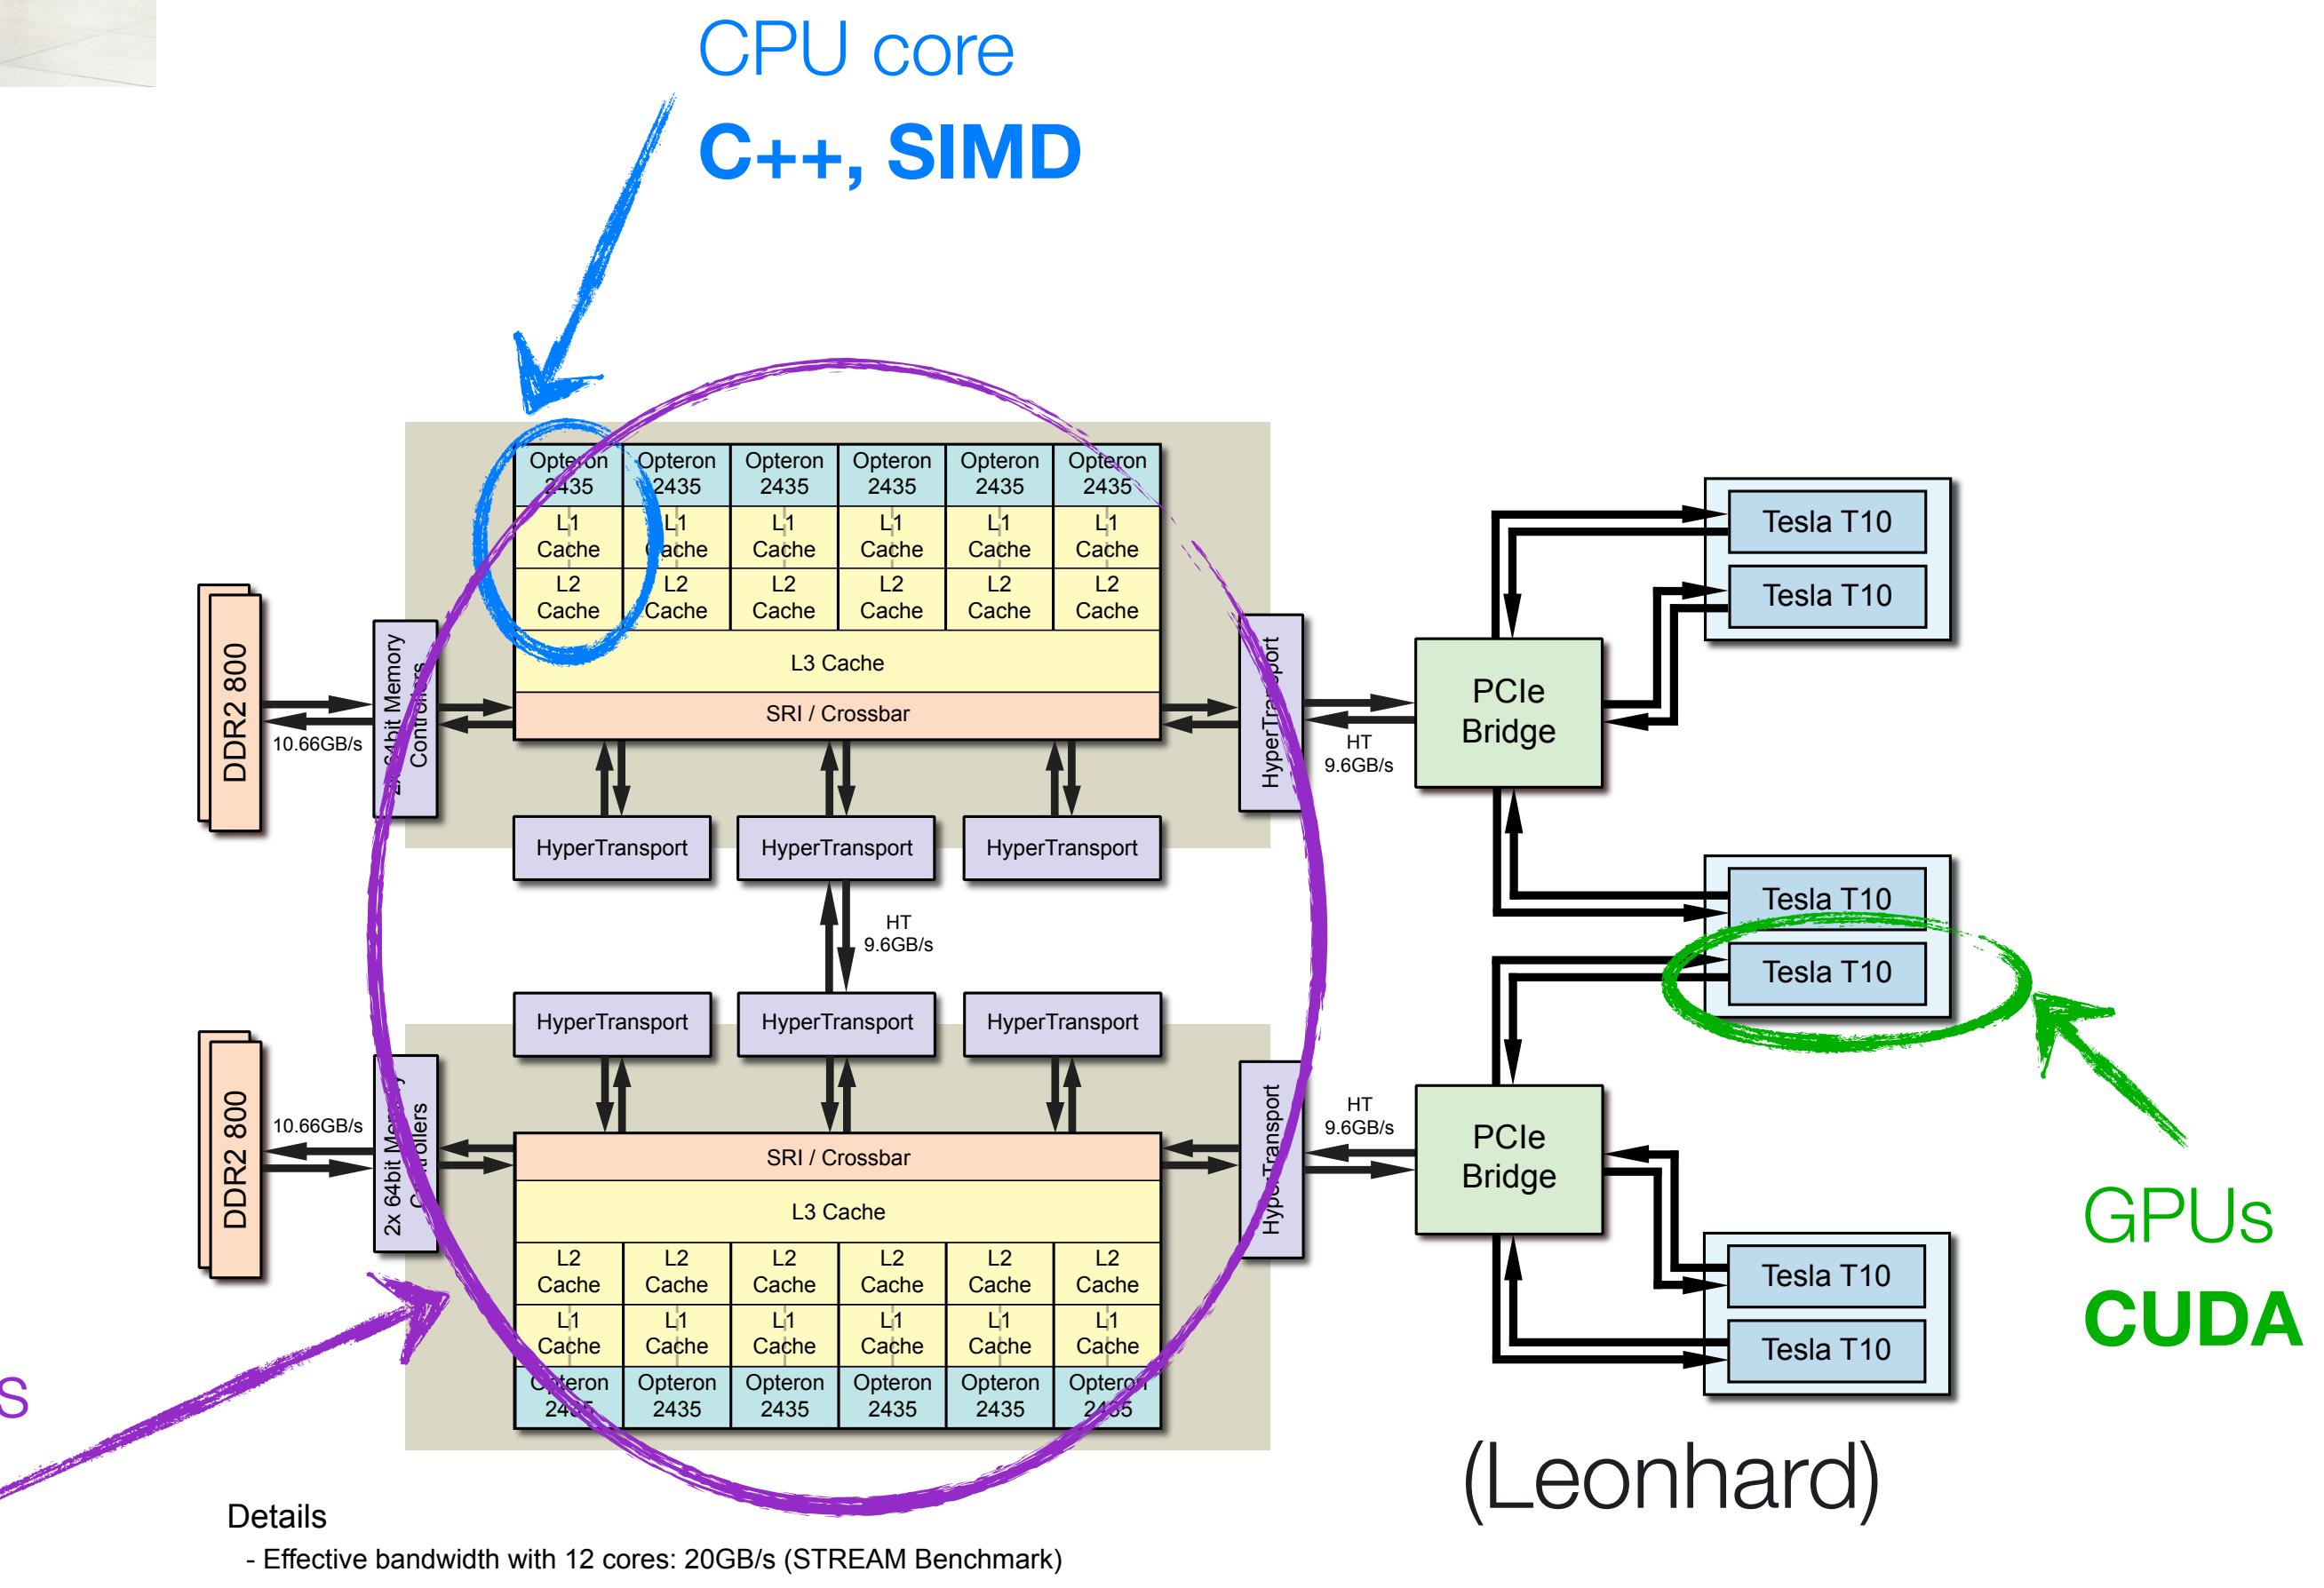
Euler



Cluster: network of nodes
Distributed memory

MPI

Node: multiple processors
Shared Memory
C++ threads, OpenMP

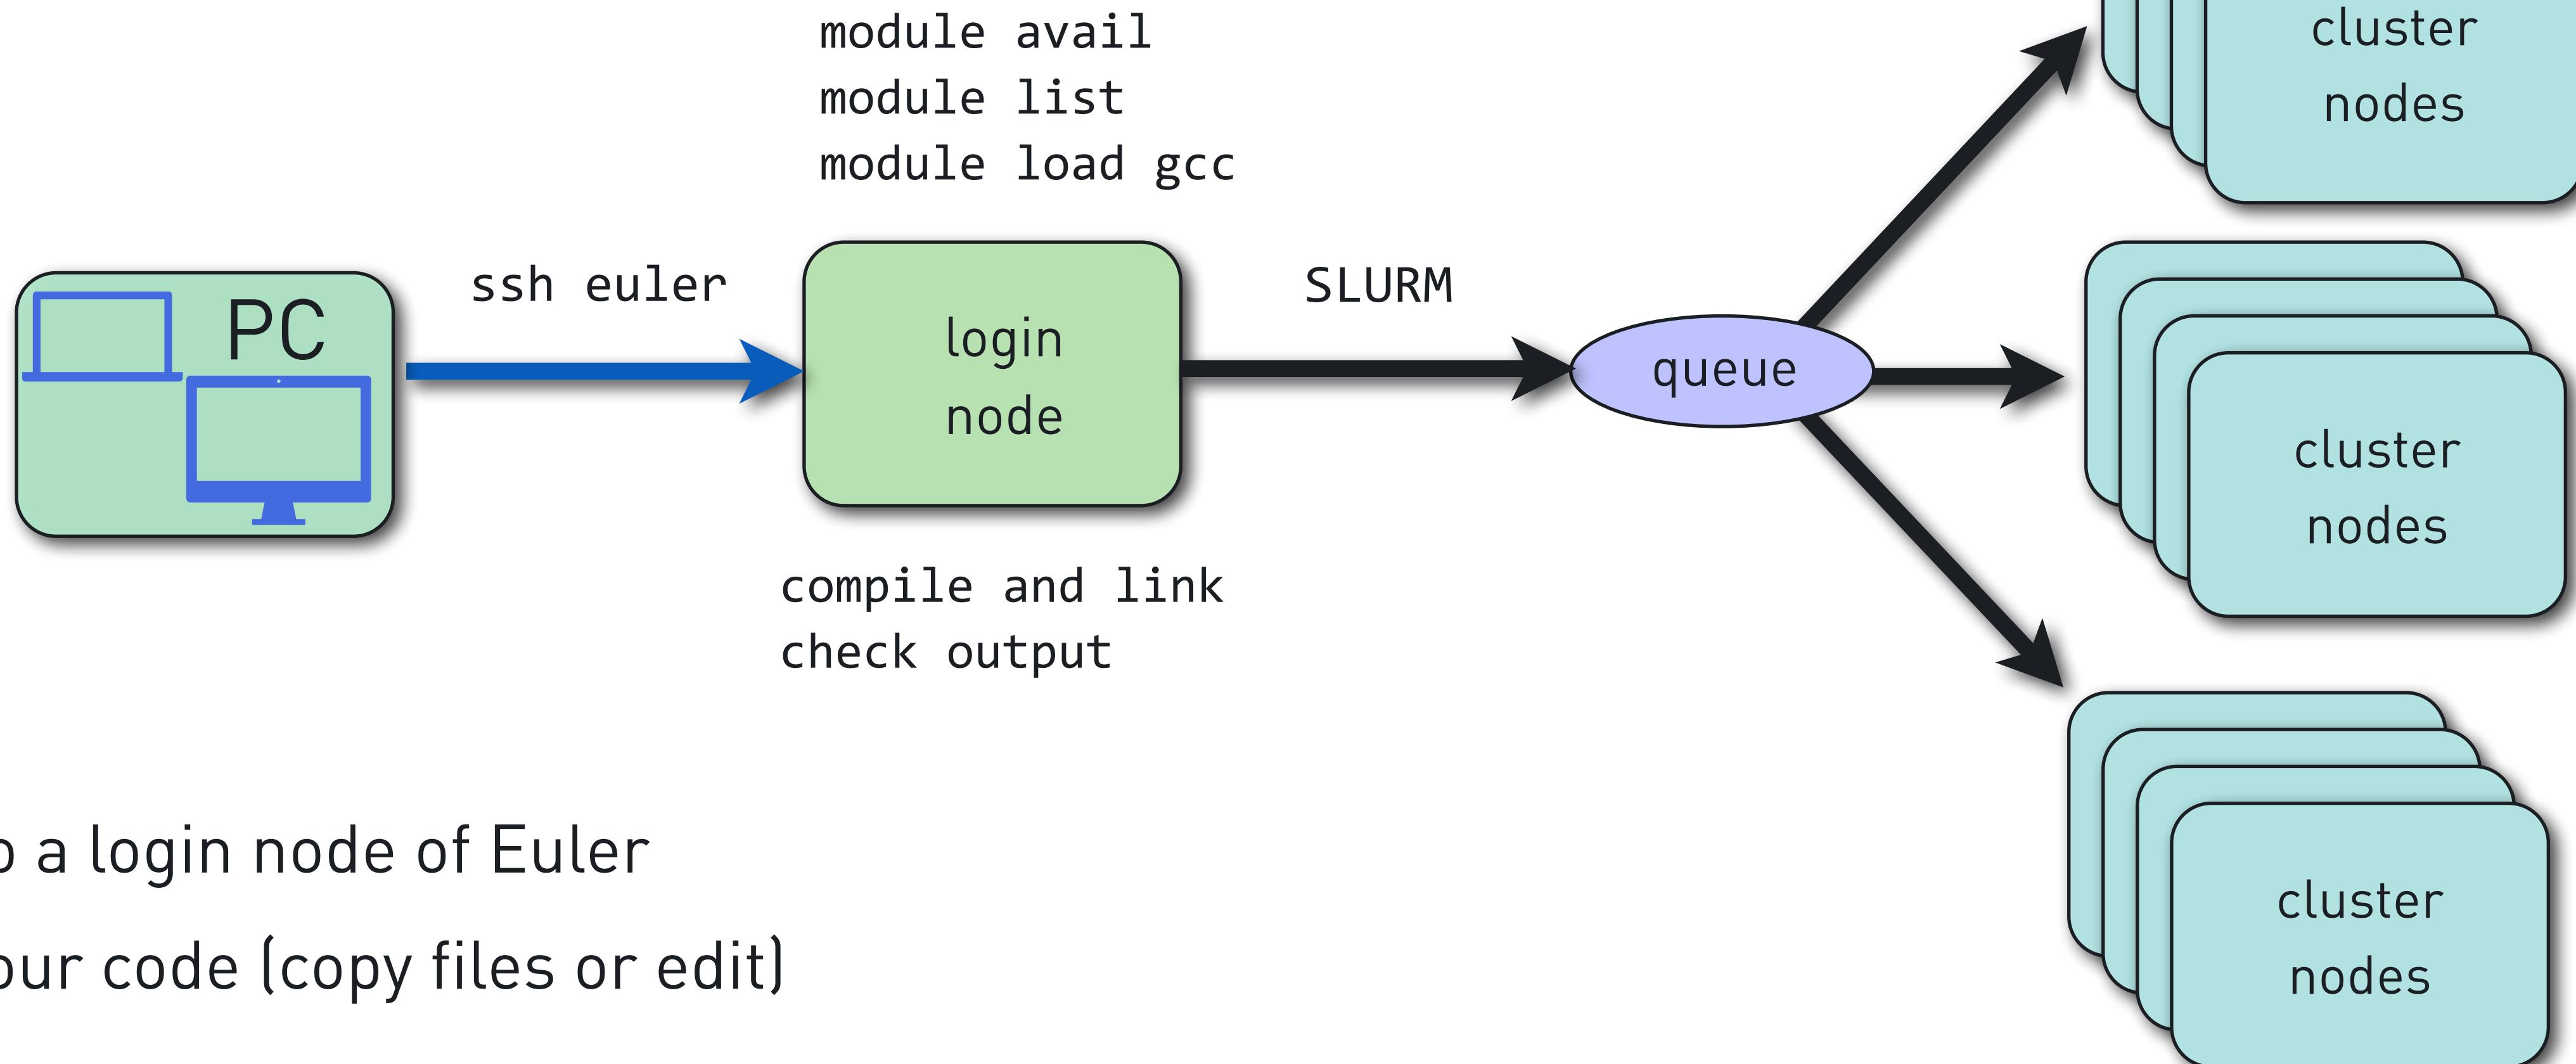


(Leonhard)

IV. Euler cluster : usage

DOCUMENTATION

https://scicomp.ethz.ch/wiki/User_documentation



STEPS

1. Connect to a login node of Euler
2. Develop your code (copy files or edit)
3. Compile your program
4. Submit a job / run your program on compute nodes
5. Check your job (status and output)

IV. Euler cluster : usage

1. Connect

- `ssh -Y <username>@euler.ethz.ch`
 - -Y: Enables trusted X11 forwarding
 - Access to one of the Euler login nodes
 - https://scicomp.ethz.ch/wiki/Accessing_the_clusters#SSH to set up SSH keys to login without typing password every time.

2. Develop

- Copy your files to Euler, e.g.
 - `scp code.tar.gz <username>@euler.ethz.ch:code.tar.gz`
 - `scp -r <Local_DIR> <username>@euler.ethz.ch:/cluster/scratch/<username>`
 - better: use Git!
- Use a text editor to write/modify your code, i.e. vi(m), emacs, nano
- OR mount your \$HOME folder locally
 - `sshfs <username>@euler.ethz.ch:/cluster/home/<username>/<local_DIR_Euler>`
 - <https://www.digitalocean.com/community/tutorials/how-to-use-sshfs-to-mount-remote-file-systems-over-ssh>

IV. Euler cluster : usage

3. Compile

- You need the appropriate programming tools and libraries to compile your code
- Just load the environment module you need
- Examples:
 - `module load gcc` (newer version of gcc)
 - `module load open_mpi` (MPI library)
 - `module list` (shows loaded modules)
 - `module avail` (what is available)
 - `module unload gcc` (unloads a module)
- Compile your code and produce the executable
- Example:
 - `g++ cputest.cpp -o cputest`
 - Use Makefiles!

IV. Euler cluster : usage

https://scicomp.ethz.ch/wiki/LSF_to_Slurm_quick_reference#Job_status
<https://user.cscs.ch/access/running/>

4. Submit your job

- login nodes are only for development
- The program must run on a compute node
- To do that, you must use the sbatch command: **sbatch jobsheet.sh**
- **jobsheet.sh**: a batch script containing options of the job
- You can also get an *interactive shell* on a compute node
srun --time=00:10:00 --pty bash

Batch script generated using Euler's Slurm Submission Line Advisor

<https://scicomp.ethz.ch/public/lsla/index2.html>

jobsheet.sh

```
#!/bin/bash

#SBATCH -n 4
#SBATCH --time=00:10:00
#SBATCH --job-name=test
#SBATCH --mem-per-cpu=512
#SBATCH --output=file.out
#SBATCH --error=file.err

srun ./hello_world_mpi
```

number of processor cores
max execution time HH:MM:SS
name of job in queue
requested memory per cpu
file containing program output
file containing standard error

run program

IV. Euler cluster : usage

https://scicomp.ethz.ch/wiki/LSF_to_Slurm_quick_reference#Job_status
<https://scicomp.ethz.ch/public/lsla/index2.html>
<https://user.csccs.ch/access/running/>

5. Check your job

- Useful commands
 - **sbatch** - submit a batch script
 - **squeue** - check the status of jobs on the system
 - **scancel** - delete one of your jobs from the queue
- Output files
 - **slurm-<JOBID>.out**: contains output of the program.
 - **slurm-<JOBID>.err**: contains any error messages (standard error). both files will be found in the directory from which you launched the job.

Simplifying our workflow

aliases

- Create an *alias* to contain the SSH command:

```
$ alias EU="ssh -Y eceva@euler.ethz.ch"
```

```
$ EU
```

- Password-less SSH:

```
$ alias EU="ssh -Y eceva@euler.ethz.ch -i ~/.ssh/euler_ssh_id"
```

```
$ EU
```

V. C++ refresher

We want to calculate and print the factorials of all numbers up to one non-negative integer number N .

The factorial of each number $n \leq N$ is calculated as:

$$n! = n \cdot (n - 1) \cdot (n - 2) \cdots 3 \cdot 2 \cdot 1 \text{ and } 0! = 1$$

Simplest
recurrent
algorithm

Hands-on c++ examples with:

- ▶ Different print, output, input utilities
- ▶ Function declarations
- ▶ Header files
- ▶ Template functions
- ▶ Class declaration

In git: tutorials/tutorial01_Intro/

```
int product = 1;
for(int n = 0; n <= N; ++n)
{
    if (n != 0)
    {
        // calculate the next factorial
        product *= n;
    }
    // output the factorial
    printf("%3d %15d\n", n, product);
}
```

C++ REFERENCE <https://en.cppreference.com/w/>

C++ TUTORIALS <https://www.tutorialspoint.com/cplusplus/>
<https://www.geeksforgeeks.org/c-plus-plus/?ref=lbp>

C++ BOOK <https://polybox.ethz.ch/index.php/s/yGjq4w0xBYsCRKT>